

VU Autonomous Racing Cars (2020S)

Lab 8: Reinforcement Learning

Team 3

Stefan Adelman


Hannes Brantner

Daniel Lukitsch

Thomas Pintaric

Overview

- Training was done on the **f1_aut** track.
- We used the **Tensorforce*** framework
 - *) <https://github.com/tensorforce/tensorforce>
 - Proximal Policy Optimization (PPO), using mostly default hyperparameters
 - Network: Two-layer CNN
 - 1:9 laser scan subsampling (120 rays)
 - Small batch size (10 episodes)
 - Randomized exploration (additional action noise)
 - Trained for 3K episodes (about 10K would be ideal).



Reward shaping

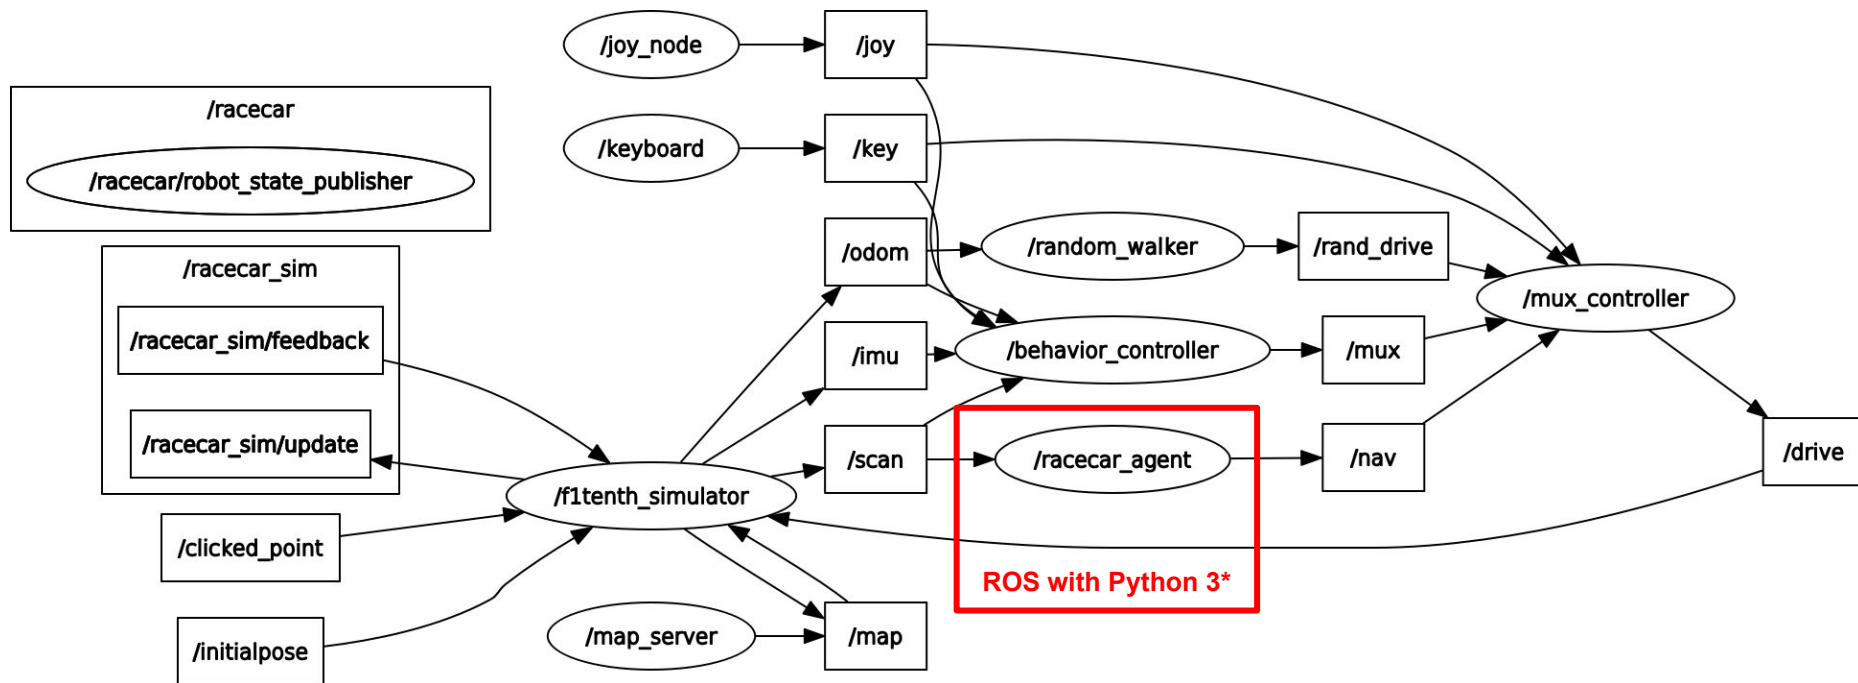
Reward is based on two **pre-computed costmaps**:
Driven distance, distance from closest obstacle.



State space exploration (after 2800 episodes):



Runtime deployment (ROS)



*) <https://gist.github.com/pintaric/d022116adfecf179adab47a196e98bce>